

# Nic-Emu

## Hyper-scalability of Network Interface Cards for Virtual Machines

Florian Dominik Freudiger

Advisor: Peter Okelmann

Chair of Computer Systems

<https://dse.in.tum.de/>



15.07.2023 – 15.11.2023

# Motivation: Firecracker

- Virtual machine monitor (VMM) managing microVMs
- Scalability:
  - Thousands of virtual machines may run on the same machine
  - Hundreds of VMs may be started within seconds
- Virtual machines need fast networking

# Existing networking solutions

- NIC passthrough:
  - High performance
  - Limited scalability
  
- NIC emulation:
  - Limited performance
  - High scalability

⇒ Hybrid solution needed

# Hybrid approach

- Switch between networking methods dynamically
  - Re-assign VMs by introducing multiplexer in-between

## ⇒ vMux

- Multiplexing in userspace
  - NIC Passthrough via Virtual Function I/O (VFIO) framework
  - NIC Emulation via behavioral models
- 
- Previous behavioral models are not easily controllable

Extensible and modular NIC behavioral model needed

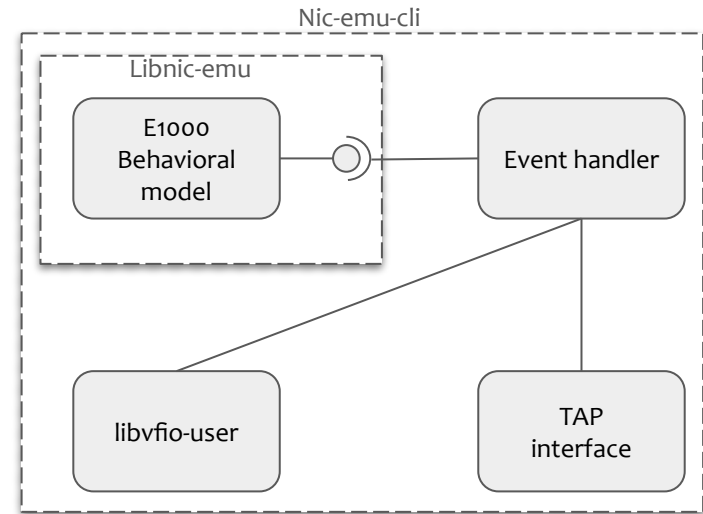
## Userspace NIC emulation

### **System design goals:**

- Driver support
- Modularity
- Performance

# System overview

- **Library: Libnic-emu**
  - Behavioral model: Intel E1000
    - Simple real NIC, cannot use paravirtualized NIC for multiplexing with real NICs
    - QEMU Reference implementation
- **Standalone: Nic-emu-cli**
  - Command-line-interface
  - Hypervisor integration: libvfiio-user
    - PCI device implementation framework
    - Supports QEMU
  - Send and receive using TAP interface




- Written in Rust
  - Since libvfio-user only exposes C headers, [libvfio-user-rs](#) wrapper has been created
- Behavioral model
  - 23/130 registers emulated
  - VLAN control, debug and statistical registers omitted
- Libnic-emu
  - Available as static library for C/C++ and as Rust crate
- Nic-emu-cli
  - Can process libvfio-user commands, packets and timers in any order using linux I/O multiplexing

- What is the performance overhead of moving emulation into userspace?
  - Comparing QEMU reference NIC emulation
- What is the performance overhead of not using paravirtualization?
  - Comparing both E1000 NICs of QEMU and Nic-emu to QEMU VirtIO



- Experimental setup (all hosts):
  - Intel Xeon Gold 5317 CPU (3.00 GHz, 12 cores)
  - 256GB DDR4 DRAM @ 2933 MT/s
- Bandwidth measurement
  - Host  $\Leftrightarrow$  Guest system
  - Via iPerf3 TCP transmission
- Latency measurement
  - Dedicated load-generator  $\Leftrightarrow$  VM Host  $\Leftrightarrow$  VM
  - Via Moongen packet generator & XDP reflector

# Bandwidth Test



Nic-emu only slower than QEMU in RX



VirtIO greatly surpasses both E1000 implementations

# Latency test

60B packets at

Low load: 10kpps

High load: 100kpps



← Lower is better

Nic-emu and QEMU E1000 perform similar on average

VirtIO remains faster

- Emulated E1000 NIC of Nic-emu performs very similar to its counterpart in QEMU in most cases
  - $\Rightarrow$  Overhead of moving NIC emulation into userspace is minor
- Para-virtualization remains much faster
  - $\Rightarrow$  Overhead of emulating real NICs needs to be overcome

**Try it out!**

<https://github.com/vmexIO/nic-emu>